

# Local Buddy: Voice + Vision + Audio on macOS (Roadmap & Starter Tasks)

## GOAL

Build a local macOS "buddy" app with voice in/out, optional vision, and on-device LLM. Start small with Discord

## PHASE 0 — BASELINE (DONE OR IN-PROGRESS)

- Git installed and linked to GitHub
- Tailscale configured
- Docker installed (optional but useful)
- Brew available

## PHASE 1 — START SMALL: DISCORD BOT (2-3 SESSIONS)

- 1) Repo setup
  - Create repo: discordbot1.0 (done)
  - Add .env with DISCORD\_TOKEN
  - npm init -y
- 2) Minimal bot
  - Dependencies:  
npm i discord.js dotenv
  - Scripts:  
npm pkg set type="module" scripts.start="node src/bot.js"
  - Health check command: "/health" → "ok"
  - Run locally: npm start
  - Containerize later with Dockerfile (optional)
- 3) Operational habits
  - One ticket per change
  - git add -A && git commit -m "feat: X" && git push
  - Keep README with run instructions

## PHASE 2 — LOCAL CLI "BUDDY" (VOICE IN/OUT) (3-5 SESSIONS)

Concept: A terminal app that listens to your mic, transcribes, sends to a local LLM, speaks the reply, and c

- 1) Choose runtime
  - Python (simplest integration) with virtualenv
  - or Node.js if you prefer JS; Python is recommended for Whisper integrations
- 2) Local LLM backend
  - Install Ollama and pull one model:  
curl -fsSL https://ollama.com/install.sh | sh  
ollama pull llama3:8b-instruct  
(Optionally: deepseek-coder:7b for code tasks)
- 3) Speech-to-text (offline)
  - Option A: faster-whisper (GPU-accelerated on Apple Silicon)  
python -m venv .venv && source .venv/bin/activate  
pip install faster-whisper sounddevice numpy  
Notes: needs mic permission in macOS
  - Option B: Vosk (lower resource)  
pip install vosk sounddevice
- 4) Text-to-speech (macOS native)
  - Use "say" command for a first version:  
say "Hello from your buddy"
  - Python wrapper:  
import subprocess; subprocess.run(["say", text])
- 5) Wire it together (Python sketch)
  - mic → faster-whisper → text
  - text → Ollama (HTTP API at http://localhost:11434)
  - response → say
  - Command safety: only run whitelisted maintenance tasks
- 6) Safe maintenance commands (whitelist)
  - system\_info: sw\_vers; uname -a
  - brew\_outdated: brew outdated
  - disk\_free: df -h /
  - memory: vm\_stat; memory\_pressure
  - Never run destructive actions without explicit confirmation
- 7) File layout

```

local-buddy/
src/
  main.py          # orchestrates loop
  stt.py           # faster-whisper integration
  tts.py           # macOS say wrapper
  llm.py           # Ollama client
  skills.py        # whitelisted commands
  .env             # config (optional)
  README.md

```

PHASE 3 – VISION (OPTIONAL, 2-4 SESSIONS)

- 1) Screen capture permissions (System Settings → Privacy & Security → Screen Recording)
- 2) OCR for reading the screen or images
  - tesseract + pytesseract (brew install tesseract; pip install pytesseract)
- 3) Basic pipeline
  - capture screenshot → OCR → summarize via LLM
- 4) Camera input (optional)
  - OpenCV for webcam frames; ask user permission first
- 5) Safety
  - Only process local images/screen on demand
  - No outbound image data without explicit permission

PHASE 4 – TURN CLI INTO A MENUBAR APP (2-4 SESSIONS)

- 1) SwiftUI wrapper (preferred) or Electron
- 2) Hotkey to start/stop listening
- 3) Status icon in menu bar
- 4) Persist settings (model choice, TTS voice, mic device)
- 5) Log viewer

PHASE 5 – DOCKERIZATION & RUNPOD (OPTIONAL)

- 1) Containerize Discord bot (already planned)
- 2) For the local buddy, keep it on Mac for mic/tts; use RunPod only for heavy LLM if needed (tunnel via Tailnet)
- 3) If using remote model:
  - Mac: stt/tts + thin client
  - RunPod: LLM server behind Tailnet
  - Configure client to talk to http://100.x.x.x:11434 or your custom port

DAY-BY-DAY STARTER TASKS (CONSERVATIVE PACING)

Day 1 (done): Git + Tailnet baseline

Day 2: Join RunPod to Tailnet; clone discordbot repo to RunPod; run npm install

Day 3: Add /health slash command to Discord bot; push; optional Dockerfile

Day 4: Start local buddy project (Python): create venv; install faster-whisper, sounddevice

Day 5: Add Ollama backend; test prompt roundtrip

Day 6: Add macOS TTS ("say"); full mic→LLM→speech loop

Day 7: Add 3 safe maintenance commands; require confirmation for any command that modifies the system

COMMAND BLOCKS (COPY WHEN READY)

A) Python venv and faster-whisper

```

python3 -m venv .venv
source .venv/bin/activate
pip install faster-whisper sounddevice numpy requests python-dotenv

```

B) Minimal Ollama client (HTTP POST)

```

import requests, os
def ask_ollama(prompt, model="llama3:8b-instruct"):
    r = requests.post("http://localhost:11434/api/generate", json={"model": model, "prompt": prompt, "stream":
    r.raise_for_status()
    return r.json().get("response", "")
print(ask_ollama("Say hello in one short sentence."))

```

C) macOS TTS helper

```

import subprocess
def speak(text: str):
    subprocess.run(["say", text])

```

D) Simple skill whitelist idea

```

ALLOWED = {
    "system_info": ["sw_vers"],
    "disk_free": ["df", "-h", "/"],
    "brew_outdated": ["brew", "outdated"]
}
# only run commands from ALLOWED; require confirmation for anything else

```

RISK GUARDRAILS

- All commands are local and explicit; never execute arbitrary shell from LLM output
- Require user confirmation ("Are you sure?") for anything that writes to disk, installs, or deletes
- Log any command run and its output to ./logs/

- Keep models local (Ollama) unless you intentionally connect to RunPod via Tailscale

#### WHAT TO BRING TO NEXT SESSION

- Confirmation that RunPod joined Tailscale
- discordbot1.0 cloned on RunPod and npm install completed
- Choice: start "local-buddy" in Python or stay on Discord bot for one more feature

#### TEACHER SUMMARY

We are leveling up in layers: Discord bot → local CLI buddy (voice) → optional vision → optional menubar app